



UNSW
SYDNEY

9. INTRACTABILITY

Serge Gaspers

serge.gaspers@unsw.edu.au

Katie Clinch

k.clinch@unsw.edu.au

Course admin: Song Fang

cs3121@cse.unsw.edu.au

School of Computer Science and Engineering, UNSW Sydney

Term 3, 2023

1. Tractability of Algorithms

2. Polynomial Reductions

3. Optimisation Problems

4. Puzzle

Definition

A (sequential) algorithm is said to be *polynomial time* if for every input it terminates in polynomially many steps in the length of the input.

This means that there exists a natural number k (independent of the input) so that the algorithm terminates in $T(n) = O(n^k)$ many steps, where n is the size of the input.

Question

What is the *length* of an input?

Answer

It is the *number of symbols* needed to describe the input precisely.

- For example, if input x is an integer, then $|x|$ can be taken to be the number of bits in the binary representation of x .
- As we will see, the definition of polynomial time computability is quite robust with respect to how we represent inputs.
- For example, we could instead define the length of an integer x as the number of digits in the *decimal* representation of x .
- This can only change the constants involved in the expression $T(n) = O(n^k)$ but not the asymptotic bound.

- If the input is a weighted graph G , then G can be described using its adjacency list: for each vertex v_i , store a list of edges incident to v_i together with their (integer) weights represented in binary.
- Alternatively, we can represent G with its adjacency matrix.
- If the input graphs are all sparse, this can unnecessarily increase the length of the representation of the graph.
- However, since we are interested only in whether the algorithm runs in polynomial time and not in the particular degree of the polynomial bounding such a run time, this does not matter.
- In fact, every precise description without artificial redundancies will do.

Definition

A *decision problem* is a problem with a YES or NO answer.

Examples include:

- “Is the input number n a prime number?”
- “Is the input graph G connected?”
- “Does the input graph G have a cycle containing all vertices of G ?”

Definition

A decision problem $A(x)$ is in class \mathbf{P} (*polynomial time*, denoted $A \in \mathbf{P}$) if there exists a polynomial time algorithm which solves it (i.e. produces the correct output for each input x).

Definition

A decision problem $A(x)$ is in class **NP** (*non-deterministic polynomial time*, denoted $A \in \mathbf{NP}$) if there exists a decision problem $B(x, y)$ such that

- 1 for every input x , $A(x)$ is true (YES) if and only if there is some y for which $B(x, y)$ is true, and
- 2 $B(x, y)$ has an algorithm running in polynomial time in the length of x *only*.

We call y a *certificate* for x .

Question

Consider the decision problem $A(x) = \text{"is the integer } x \text{ not prime"}$.
Is $A \in \mathbf{NP}$?

Answer

- We need to find a problem $B(x, y)$ such that $A(x)$ is true if and only if there is some y for which $B(x, y)$ is true.
- The natural choice is $B(x, y) = \text{"}x \text{ is divisible by } y\text{"}$.
- $B(x, y)$ indeed has an algorithm running in time polynomial in the length of x only.

Question

Is $A \in \mathbf{P}$?

Answer

- Also yes! But this is not at all straightforward. This is a famous and unexpected result, proved in 2002 by the Indian computer scientists Agrawal, Kayal and Saxena.
- The AKS algorithm provides a *deterministic, polynomial time* procedure for testing whether an integer x is prime.

The length of the input for primality testing is $O(\log x)$.

Comparing some well-known algorithms for primality testing:

- The naïve algorithm tests all possible factors up to the square root, so it runs in $O(\sqrt{x})$ and is therefore not a polynomial time algorithm.
- The Miller-Rabin algorithm runs in time proportional to $O(\log^3 x)$ but determines only *probable primes*.
- The original AKS algorithm runs in $\tilde{O}(\log^{12} x)$, and newer versions run in $\tilde{O}(\log^6 x)$.
- However, the AKS algorithm is rarely used in practice; tests using elliptic curves are much faster.

Vertex Cover

Input: a graph G and an integer k .

Question: “Is there a subset U of at most k vertices of G (called a vertex cover of G) such that each edge has at least one endpoint belonging to U .”

- Clearly, given a subset of vertices U we can determine in polynomial time whether U is a vertex cover of G with at most k elements.
- So Vertex Cover is in **NP**.

Satisfiability (SAT)

Input: a propositional formula in conjunctive normal form (CNF) $C_1 \wedge C_2 \wedge \dots \wedge C_m$ where each clause C_i is a disjunction of propositional variables or their negations; for example

$$(P_1 \vee \neg P_2 \vee P_3 \vee \neg P_5) \wedge (P_2 \vee P_3 \vee \neg P_5 \vee \neg P_6) \wedge (\neg P_3 \vee \neg P_4 \vee P_5)$$

Question: “Is there an assignment to the propositional variables which makes the formula true?”

- Clearly, given an assignment to the propositional variables one can determine in polynomial time whether the formula is true for this assignment.
- So SAT is in **NP**.

Question

Is SAT in the class **P**?

(Partial) Answer

If each clause C_i involves exactly two variables (2SAT), then yes!

In this case, it can be solved in linear time using strongly connected components and topological sort.

Another special case of interest is when each clause involves exactly *three* variables (3SAT). This will be fundamental in our study of NP-complete and NP-hard problems.

Question

Is it the case that *every* problem in **NP** is also in **P**?

- For example, is there a polynomial time algorithm to solve the SAT problem?
- The existence of such an algorithm would mean that finding out whether a propositional formula evaluates true in any of the 2^n assignment to its n variables is actually not much harder than simply checking one of these cases.
- Intuitively, this should not be the case; determining if such a case exists should be a harder problem than simply checking a particular case.

- However, so far no one has been able to prove (or disprove) this, despite decades of effort by very many very famous people!
- The conjecture that **NP** is a strictly larger class of decision problems than **P** is known as the “**P** \neq **NP**” hypothesis, and it is widely considered to be one of the hardest open problems in mathematics.

1. Tractability of Algorithms

2. Polynomial Reductions

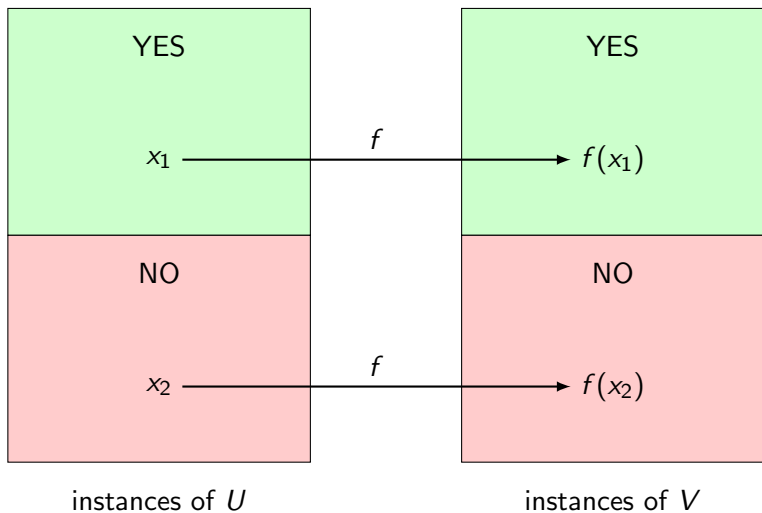
3. Optimisation Problems

4. Puzzle

Definition

Let U and V be two decision problems. We say that U is polynomially reducible to V if and only if there exists a function $f(x)$ such that:

- 1 $f(x)$ maps instances of U to instances of V .
- 2 f maps YES instances of U to YES instances of V and NO instances of U to NO instances of V ,
i.e. $U(x)$ is YES if and only if $V(f(x))$ is YES.
- 3 $f(x)$ is computable by a polynomial time algorithm.



Definition

The *contrapositive* of the implication $p \implies q$ is $\neg q \implies \neg p$.

Example

“Students who enjoy puzzles look forward to the end of each Algorithms lecture” is logically equivalent to “Students who dread the end of each Algorithms lecture don’t enjoy puzzles”.

Note

Instead of proving that if x is a NO instance, then $f(x)$ is a NO instance, we often prove the equivalent statement that if $f(x)$ is a YES instance, it must have been mapped from a YES instance x .

Claim

Every instance of SAT is polynomially reducible to an instance of 3SAT.

Proof Outline

We introduce more propositional variables and replace every clause by a conjunction of several clauses.

For example, we replace the clause

$$\underbrace{P_1 \vee \neg P_2 \vee \neg P_3}_{\text{clause 1}} \vee \underbrace{P_4 \vee \neg P_5 \vee P_6}_{\text{clause 2}} \quad (1)$$

with the following conjunction of “chained” 3-clauses with new propositional variables Q_1, Q_2, Q_3 :

$$\begin{aligned} & (\underbrace{P_1 \vee \neg P_2 \vee Q_1}_{\text{clause 1}}) \wedge (\neg Q_1 \vee \underbrace{\neg P_3 \vee Q_2}_{\text{clause 2}}) \\ & \wedge (\neg Q_2 \vee \underbrace{P_4 \vee Q_3}_{\text{clause 3}}) \wedge (\neg Q_3 \vee \underbrace{\neg P_5 \vee P_6}_{\text{clause 4}}) \end{aligned} \quad (2)$$

Easy to verify that if an evaluation of the P_i makes (1) true, then the corresponding evaluation of the Q_j also makes (2) true and vice versa: every evaluation which makes (2) true also makes (1) true. Clearly, (2) can be obtained from (1) using a simple polynomial time algorithm.

Theorem

Every NP problem is polynomially reducible to the SAT problem.

This means that for every NP decision problem $U(x)$ there exists a polynomial time computable function $f(x)$ such that:

- 1 for every instance x of U , $f(x)$ produces a propositional formula Φ_x ;
- 2 $U(x)$ is true if and only if Φ_x is satisfiable.

Definition

An NP decision problem U is *NP-complete* ($U \in \mathbf{NP-C}$) if every other NP problem is polynomially reducible to U .

- Thus, Cook's Theorem says that SAT is NP-complete.
- NP-complete problems are in a sense universal: if we had an algorithm which solves any NP-complete problem U , then we could also solve every other NP problem as follows.
- A solution of an instance x of any other NP problem V could simply be obtained by:
 - 1 computing in polynomial time the reduction $f(x)$ of V to U ,
 - 2 then running the algorithm that solves U on instance $f(x)$.

- So NP-complete problems are the hardest NP problems - a polynomial time algorithm for solving an NP-complete problem would make every other NP problem also solvable in polynomial time.
- But if $\mathbf{P} \neq \mathbf{NP}$ (as is commonly hypothesised), then there cannot be any polynomial time algorithms for solving an NP-complete problem, not even an algorithm that runs in time $O(n^{1,000,000})$.

- Maybe SAT is not that important - why should we care about satisfiability of propositional formulas?
- Maybe NP-complete problems only have theoretical significance and no practical relevance?
- Unfortunately, this could not be further from the truth!
- A vast number of practically important decision problems are NP-complete!

Traveling Salesman Problem

Input:

- 1 a map, i.e., a weighted directed graph with:
 - vertices representing locations
 - edges representing roads between pairs of locations
 - edge weights representing the lengths of these roads;
- 2 a number L .

Question: Is there a tour along the edges which visits each location (i.e., vertex) exactly once and returns to the starting location, with total length at most L ?

Think of a mailman who has to deliver mail to several addresses and then return to the post office. Can he do it while traveling less than L kilometres in total?

Register Allocation Problem

Input:

- 1 an undirected unweighted graph G with:
 - vertices representing program variables
 - edges representing pairs of variables which are both needed at the same step of program execution;
- 2 the number of registers K of the processor.

Question: is it possible to assign variables to registers so that no edge has both vertices assigned to the same register?

In graph theoretic terms: is it possible to color the vertices of a graph G with at most K colors so that no edge has both vertices of the same color?

Vertex Cover Problem

Input:

- 1 an undirected unweighted graph G with vertices and edges;
- 2 a number k .

Question: it it possible to choose k vertices so that every edge is incident to at least one of the chosen vertices?

Set Cover Problem

Input:

- 1 a number of items n ;
- 2 a number of bundles m such that
 - each bundle contains of a subset of the items
 - each item appears in at least one bundle;
- 3 a number k .

Question: it it possible to choose k bundles which together contain all n items?

This problem can be extended by assigning a price to each bundle, and asking whether satisfactory bundles can be chosen within a budget b .

- We will see that many other practically important problems are also NP-complete.
- Be careful though: sometimes the distinction between a problem in **P** and a problem in **NP-C** can be subtle!

Problems in **P**:

- Given a graph G and two vertices s and t , is there a path from s to t of length *at most* K ?
- Given a propositional formula in CNF form such that every clause has at most *two* propositional variables, does the formula have a satisfying assignment? (2SAT)
- Given a graph G , does G have a tour where every *edge* is traversed exactly once? (Euler tour)

Problems in **NP-C**:

- Given a graph G and two vertices s and t , is there a simple path from s to t of length *at least* K ?
- Given a propositional formula in CNF form such that every clause has at most *three* propositional variables, does the formula have a satisfying assignment? (3SAT)
- Given a graph G , does G have a tour where every *vertex* is visited exactly once? (Hamiltonian cycle)

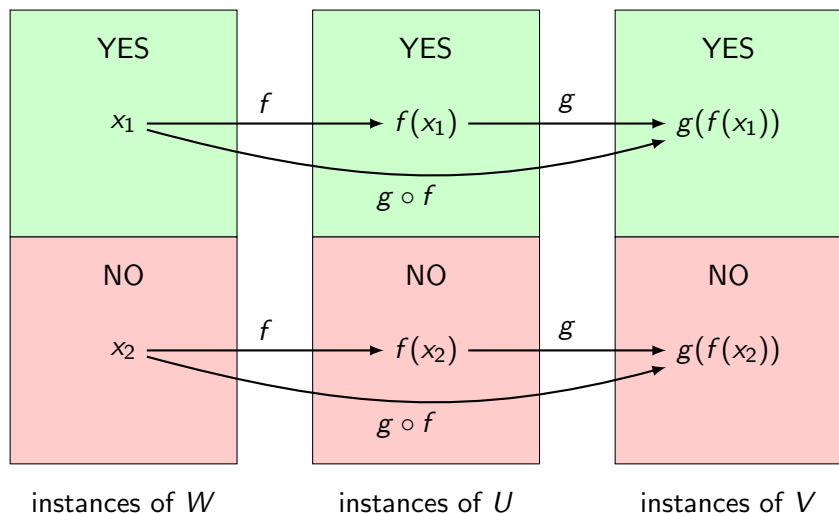
Taking for granted that SAT is NP-complete, how do we prove NP-completeness of another NP problem?

Theorem

Let U be an NP-complete problem, and let V be another NP problem. If U is polynomially reducible to V , then V is also NP-complete.

Proof

- Let $g(x)$ be a polynomial reduction of U to V , and let W be any other NP problem.
- Since U is NP-complete, there exists a polynomial reduction $f(x)$ of W to U .
- We will now prove that $(g \circ f)(x)$ is a polynomial reduction of W to V .



Proof (continued)

We first claim that $(g \circ f)(x)$ is a reduction from W to V .

- 1 Since f is a reduction from W to U , $W(x)$ is YES iff $U(f(x))$ is YES.
- 2 Since g is a reduction from U to V , $U(f(x))$ is YES iff $V(g(f(x)))$ is YES.

Thus $W(x)$ is YES iff $V(g(f(x)))$ is YES, i.e., $(g \circ f)(x)$ is a reduction of W to V .

Proof (continued)

- Since $f(x)$ is the output of a polynomial time computable function, the length $|f(x)|$ of the output $f(x)$ can be at most a polynomial in $|x|$, i.e., for some polynomial (with positive coefficients) P we have $|f(x)| \leq P(|x|)$.
- Since $g(y)$ is polynomial time computable as well, there exists a polynomial Q such that for every input y , computation of $g(y)$ terminates after at most $Q(|y|)$ many steps.
- Thus, the computation of $(g \circ f)(x)$ terminates in at most:
 - $P(|x|)$ many steps, for the computation of $f(x)$, plus
 - $Q(|f(x)|) \leq Q(P(|x|))$ many steps, for the computation of $g(y)$ (where $y = f(x)$).
- In total, the computation of $(g \circ f)(x)$ terminates in at most $P(|x|) + Q(P(|x|))$ many steps, which is polynomial in $|x|$.

Proof (continued)

- Therefore $(g \circ f)(x)$ is a polynomial reduction from W to V .
- But W could be any NP problem!
- We have now proven that any NP problem is polynomially reducible to the NP problem V , i.e., V is NP-complete.

Problem

Prove that Vertex Cover (VC) is NP-complete by finding a polynomial time reduction from 3SAT to VC.

Outline

We will map each instance Φ of 3SAT to a corresponding instance $f(\Phi) = (G, k)$ of VC in polynomial time, and prove that:

- 1 if Φ is a YES instance of 3SAT, then $f(\Phi)$ is a YES instance of VC, and
- 2 if $f(\Phi)$ is a YES instance of VC, then Φ is a YES instance of 3SAT.

Note that this uses the earlier mentioned contrapositive.

Construction

Given an instance of 3SAT, start with an empty graph for VC:


- 1 for each propositional variable P_k add a pair of new adjacent “literal” vertices labeled P_k and $\neg P_k$;
- 2 for each clause C_i , add a triangle with three vertices v_1^i, v_2^i, v_3^i and three edges connecting these vertices; each of these vertices represents a literal (i.e., a variable or its negation) occurring in C_i ;
- 3 for each of these vertices v_j^i , if it represents the literal L_k (which is P_k or $\neg P_k$) then make it adjacent to the vertex labeled L_k .

Example

Consider the propositional formula

$$(P_1 \vee \neg P_2 \vee P_3) \wedge (\neg P_1 \vee P_2 \vee P_3) \wedge (\neg P_1 \vee P_2 \vee \neg P_3).$$

The corresponding graph for Vertex Cover is:



Claim

An instance of 3SAT consisting of m clauses and n propositional variables is satisfiable if and only if the corresponding graph has a vertex cover of size at most $2m + n$.

Proof

Assume there is a vertex cover with at most $2m + n$ vertices chosen. Then

- 1 each triangle must have at least two vertices chosen, and
- 2 each pair of literal vertices must have at least one of its vertices chosen.


This is in total $2m + n$ vertices; thus each triangle must have *exactly* two vertices chosen and each pair of literal vertices must have *exactly* one of its vertices chosen.

Proof (continued)

Recall the example

$$(P_1 \vee \neg P_2 \vee P_3) \wedge (\neg P_1 \vee P_2 \vee P_3) \wedge (\neg P_1 \vee P_2 \vee \neg P_3)$$

and the corresponding VC instance.



Proof (continued)

- For each propositional variable P_i , set it to true if the vertex labelled P_i is in the vertex cover and to false if $\neg P_i$ is in the vertex cover.
- Now, in each triangle, for the vertex that is not in the vertex cover, all its neighbors must be in the vertex cover, in particular the neighbor that represents a literal that has been set to true; this guarantees that the clause corresponding to each triangle is true.

Proof (continued)

- For the reverse direction, assume that the formula has an assignment of the variables which makes it true.
- For each propositional variable P_k , if the assignment sets P_k to true, then add the vertex labelled P_k to the vertex cover; otherwise add the vertex labelled $\neg P_k$.
- For each clause, select one of the literals that is true under the assignment, and add the two triangle vertices that correspond to the other literals to the vertex cover.
- In this way we add exactly $2m + n$ vertices of the graph to the vertex cover and every edge between a literal pair and a triangle is covered.

1. Tractability of Algorithms

2. Polynomial Reductions

3. Optimisation Problems

4. Puzzle

- Let A be a problem and suppose we have a “black box” device which for every input x instantaneously computes $A(x)$. We also refer to such black boxes as **oracles**.
- We consider algorithms which are *polynomial time in A* . This means algorithms which run in polynomial time in the length of the input and which, besides the usual computational steps, can also use the oracle.

Definition

We say that a problem A is *NP-hard* ($A \in \mathbf{NP-H}$) if every NP problem is polynomial time in A , i.e., if we can solve every NP problem U using a polynomial time algorithm which can also use a black box to solve any instance of A .

Note that we do NOT require that A be an NP problem; we even do not require it to be a decision problem - it can also be an optimisation problem.

Traveling Salesman Optimisation Problem

Instance:

- 1 a map, i.e., a weighted graph with:
 - vertices representing locations
 - edges representing roads between pairs of locations
 - edge weights representing the lengths of these roads;

Problem: find a tour along the edges which visits each location (i.e., vertex) exactly once and returns to the starting location, with *minimal* total length?

Think of a mailman who has to deliver mail to several addresses and then return to the post office. How can he do it while traveling the minimum total distance?

- The Traveling Salesman Optimisation Problem is clearly NP-hard: using a “black box” for solving it, we can solve the Traveling Salesman Decision problem.
- Traveling Salesman Decision problem: given a weighted graph G and a number L is there a tour containing all vertices of the graph and whose length is at most L .
- We simply invoke the black box for the Traveling Salesman Optimisation Problem, which gives us the length of a shortest tour, and compare this length with L .
- Since the Traveling Salesman Decision Problem is NP-complete, all other NP problems are polynomial time reducible to it.
- Therefore every other NP problem is solvable using a “black box” for the Traveling Salesman Optimisation Problem.

- It is important to be able to figure out if a problem at hand is NP-hard in order to know that one has to abandon trying to come up with a feasible (i.e., polynomial time) solution.
- So what do we do when we encounter an NP-hard problem?
- If this problem is an optimisation problem, one approach is to try to solve it in an approximate sense by finding a solution which might not be optimal, but it is reasonably close to an optimal solution.
- For example, in the case of the Traveling Salesman Optimisation Problem we might look for a tour which is at most twice the length of the shortest possible tour.

- Thus, for a practical problem which appears to be intractable, the strategy would be:
 - prove that the problem is indeed NP-hard, to justify not trying solving the problem exactly in polynomial time;
 - look for an approximation algorithm which provides a feasible sub-optimal solution that it is not too far from optimal.

Algorithm

- 1 Pick an arbitrary edge and add both endpoints to the vertex cover.
- 2 Remove the two vertices and all the edges that are now covered. All remaining edges still need to be covered.
- 3 Continue picking edges until no edges are left.

- This certainly produces a vertex cover, because edges are removed only if they are covered, and we perform this procedure until no edges are left.
- The size of this vertex cover is equal to twice the number of edges covered by both endpoints.
- But the minimal vertex cover must include at least one vertex of each such edge.
- Thus we have produced a vertex cover of size at most twice the size of the minimum vertex cover.

Problem

Input: a complete weighted graph G with weights $d(i, j)$ of edges (to be interpreted as distances) satisfying the “triangle inequality”: for any three vertices i, j, k , we have

$$d(i, j) + d(j, k) \geq d(i, k).$$


Claim

MTSP has an approximation algorithm producing a tour of total length at most twice the length of the optimal (i.e., shortest) length tour, which we will denote by opt .

Algorithm


Find a minimum spanning tree T of G . Since the optimal tour with one of its edges e removed is a spanning tree, we have that the total weight of T satisfies $w(T) \leq \text{opt} - w(e) \leq \text{opt}$.

If we do a depth first traversal of the tree, we will travel a total distance of $2w(T) \leq 2\text{opt}$.



Algorithm (continued)

We now take shortcuts to avoid visiting vertices more than once; because of the triangle inequality, this operation does not increase the length of the tour.



- All NP-complete problems are equally difficult, because any of them is polynomially reducible to any other.
- However, the related approximation problems can be very different!
- For example, we have seen that some of these optimisation problems allow us to get within a constant factor of the optimal answer.
 - Vertex Cover permits an approximation which produces a cover at most twice as large as the minimum vertex cover.
 - Metric TSP permits an approximation which produces a tour at most twice as long as the shortest tour.

- On the other hand, the general Traveling Salesman Problem does not allow a constant-factor polynomial-time approximation algorithm: if $\mathbf{P} \neq \mathbf{NP}$, then for no $K > 1$ can there be a polynomial time algorithm which for every instance produces a tour which is at most K times the length of the shortest tour!
- To prove this, we show that if for some $K > 0$ there was indeed a polynomial time algorithm producing a tour which is at most K times the length of the shortest tour, then we could obtain a polynomial time algorithm which solves the NP-complete Hamiltonian Cycle problem.
- This is the problem of determining for a graph G whether G contains a cycle visiting all vertices exactly once.

- Let G be an arbitrary unweighted graph with n vertices.
- We turn this graph into a complete weighted graph G^* by setting the weights of all existing edges to 1, and then adding edges of weight $K \cdot n + 1$ between the remaining pairs of vertices.
- If an approximation algorithm for TSP exists, it produces a tour of all vertices with total length at most $K \cdot \text{opt}$, where opt is the length of the optimal tour through G^* .

- If the original graph G has a Hamiltonian cycle, then G^* has a tour consisting of edges already in G and of weights equal to 1, so such a tour has length of exactly n .
- Otherwise, if G does not have a Hamiltonian cycle, then the optimal tour through G^* must contain at least one added edge of length $K \cdot n + 1$, so

$$opt \geq (K \cdot n + 1) + (n - 1) \cdot 1 > K \cdot n.$$

- Thus, our approximate TSP algorithm either returns:
 - a tour of length at most $K \cdot n$, indicating that G has a Hamiltonian cycle, or
 - a tour of length greater than $K \cdot n$, indicating that G does not have a Hamiltonian cycle.
- If this approximation algorithm runs in polynomial time, we now have a polynomial time decision procedure for determining whether G has a Hamiltonian cycle!
- This can only be the case if **P** = **NP**.

- Exact exponential time algorithms. There is an algorithm solving Vertex Cover in $O(1.1970^n)$ time.
- Fixed parameter algorithms. There is an algorithm solving Vertex Cover in $O(1.2738^k + kn)$ time.
- Heuristics. The COVER heuristic (COVER Edges Randomly) finds a smaller vertex cover than state-of-the-art heuristics on a suite of hard benchmark instances.
- Reductions and solvers. There is a polynomial-time reduction from Vertex Cover to SAT, and a well-engineered SAT solver, which solves Vertex Cover fast in practice.
- Restricting the inputs. Vertex Cover can be solved in polynomial time on bipartite graphs, trees, interval graphs, etc.
- Quantum algorithms. There is a bounded-error quantum algorithm solving Vertex Cover in $O(1.0941^n)$ time.

1. Tractability of Algorithms
2. Polynomial Reductions
3. Optimisation Problems
4. Puzzle

Problem

You are given a coin, but you are not guaranteed that it is a fair coin. It may be biased towards either heads or tails. Use this coin to simulate a fair coin.

Hint

Try tossing the biased coin more than once!



That's All, Folks!!